[57] R. J. Williams, "Simple statistical gradient-following algorithms for connectionist reinforcement learning," *Machine Learning*, vol. 8, no. 3, pp. 229–256, 1992.

[58] T. Shen, T. Zhou, G. Long, J. Jiang, S. Wang, and C. Zhang, "Reinforced self-attention network: a hybrid of hard and soft attention for sequence modeling," in *27th International Joint Conference on Artificial Intelligence (IJCAI 2018)*. IJCAI, 2018, pp. 4345–4352.

[59] M. Malinowski, C. Doersch, A. Santoro, and P. Battaglia, "Learning visual question answering by bootstrapping hard attention," in *2018 European Conference on Computer Vision (ECCV 2018)*, 2018.

[60] Y. Liu, W. Wang, Y. Hu, J. Hao, X. Chen, and Y. Gao, "Multi-agent game abstraction via graph attention neural network," in *34th AAAI Conference on Artificial Intelligence (AAAI 2020)*, vol. 34, no. 05. AAAI Press, 2020, pp. 7211–7218.

[61] S. Seo, J. Huang, H. Yang, and Y. Liu, "Interpretable convolutional neural networks with dual local and global attention for review rating prediction," in *11th ACM Conference on Recommender Systems (RecSys 2017)*. ACM, 2017, pp. 297–305.

[62] J. Wang, C. Sun, S. Li, X. Liu, L. Si, M. Zhang, and G. Zhou, "Aspect sentiment classification towards question-answering with reinforced bidirectional attention network," in *57th Annual Meeting of the Association for Computational Linguistics (ACL 2019)*. ACL, 2019, pp. 3548–3557.

[63] M. Jiang, C. Li, J. Kong, Z. Teng, and D. Zhuang, "Cross-level reinforced attention network for person re-identification," *Journal of Visual Communication and Image Representation*, vol. 69, p. 102775, 2020.

[64] T. Shen, T. Zhou, G. Long, J. Jiang, S. Pan, and C. Zhang, "DiSAN: Directional self-attention network for RNN/CNN-free language understanding," in *32nd AAAI Conference on Artificial Intelligence (AAAI 2018)*. AAAI Press, 2018, pp. 5446–5455.

[65] O. Arshad, I. Gallo, S. Nawaz, and A. Calefati, "Aiding intra-text representations with visual context for multimodal named entity recognition," in *2019 International Conference on Document Analysis and Recognition (ICDAR 2019)*. IEEE, 2019, pp. 337–342.

[66] W. Wu, X. Sun, and H. Wang, "Question condensing networks for answer selection in community question answering," in *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL 2018)*. ACL, 2018, pp. 1746–1755.

[67] O. Oktay, J. Schlemper, L. L. Folgoc, M. Lee, M. Heinrich, K. Misawa, K. Mori, S. McDonagh, N. Y. Hammerla, B. Kainz, B. Glocker, and D. Rueckert, "Attention U-Net: Learning where to look for the pancreas," in *1st Medical Imaging with Deep Learning Conference (MIDL 2018)*, 2018.

[68] R. Tan, J. Sun, B. Su, and G. Liu, "Extending the transformer with context and multi-dimensional mechanism for dialogue response generation," in *8th International Conference on Natural Language Processing and Chinese Computing (NLPCC 2019)*, ser. LNCS, J. Tang, M.-Y. Kan, D. Zhao, S. Li, and H. Zan, Eds., vol. 11839. Springer, 2019, pp. 189–199.

[69] L. Chen, B. Lv, C. Wang, S. Zhu, B. Tan, and K. Yu, "Schema-guided multi-domain dialogue state tracking with graph attention neural networks," in *34th AAAI Conference on Artificial Intelligence (AAAI 2020)*, vol. 34, no. 05. AAAI Press, 2020, pp. 7521–7528.

[70] H. Wang, G. Liu, A. Liu, Z. Li, and K. Zheng, "Dmran: A hierarchical fine-grained attention-based network for recommendation," in *28th International Joint Conference on Artificial Intelligence (IJCAI 2019)*.

[71] Z. Lin, M. Feng, C. N. d. Santos, M. Yu, B. Xiang, B. Zhou, and Y. Bengio, "A structured self-attentive sentence embedding," in *5th International Conference on Learning Representations (ICLR 2017)*, 2017.

[72] J. L. Ba, J. R. Kiros, and G. E. Hinton, "Layer normalization," *arXiv:1607.06450*, 2016.

[73] H. Zhao, J. Jia, and V. Koltun, "Exploring self-attention for image recognition," in *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2020)*, 2020, pp. 10 076–10 085.

[74] H. Zhang, I. Goodfellow, D. Metaxas, and A. Odena, "Self-attention generative adversarial networks," in *36th International Conference on Machine Learning (ICML 2019)*, ser. Proceedings of Machine Learning Research, K. Chaudhuri and R. Salakhutdinov, Eds., vol. 97. PMLR, 2019, pp. 7354–7363.

[75] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial nets," in *27th Annual Conference on Neural Information Processing Systems (NIPS 2014)*. Curran Associates, Inc., 2014, pp. 2672–2680.

[76] A. Sinha and J. Dolz, "Multi-scale self-guided attention for medical image segmentation," *IEEE Journal of Biomedical and Health Informatics*, vol. 25, no. 1, pp. 121–130, 2021.

[77] J. Fajtl, H. S. Sokeh, V. Argyriou, D. Monekosso, and P. Remagnino, "Summarizing videos with attention," in *2018 Asian Conference on Computer Vision (ACCV 2018)*, ser. LNCS, vol. 11367. Springer, 2018, pp. 39–54.

[78] J. Salazar, K. Kirchhoff, and Z. Huang, "Self-attention networks for connectionist temporal classification in speech recognition," in *2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP 2019)*. IEEE, 2019, pp. 7115–7119.

[79] T. Afouras, J. S. Chung, A. Senior, O. Vinyals, and A. Zisserman, "Deep audio-visual speech recognition," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, pp. 1–1, 2018.

[80] S. Zhang, Y. Tay, L. Yao, and A. Sun, "Next item recommendation with self-attention," *arXiv preprint arXiv:1808.06414*, 2018.

[81] G. Letarte, F. Paradis, P. Giguère, and F. Laviolette, "Importance of self-attention for sentiment analysis," in *2018 Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP (BlackboxNLP 2018)*. ACL, 2018, pp. 267–275.

[82] A. Sankar, Y. Wu, L. Gou, W. Zhang, and H. Yang, "Dysat: Deep neural representation learning on dynamic graphs via self-attention networks," in *13th International Conference on Web Search and Data Mining (WSDM 2020)*, 2020, pp. 519–527.

[83] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, "Graph attention networks," in *5th International Conference on Learning Representations (ICLR 2017)*, 2017.

[84] S. Iida, R. Kimura, H. Cui, P.-H. Hung, T. Utsuro, and M. Nagata, "Attention over heads: A multi-hop attention for neural machine translation," in *57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop (ACL-SRW 2019)*. ACL, 2019, pp. 217–222.

[85] N. K. Tran and C. Niedereée, "Multihop attention networks for question answer matching," in *41st ACM SIGIR International Conference on Research & Development in Information Retrieval (SIGIR 2018)*. ACM, 2018, pp. 325–334.

[86] Y. Gong and S. R. Bowman, "Ruminating reader: Reasoning with gated multi-hop attention," in *5th International Conference on Learning Representation (ICLR 2017)*, 2017.

[87] S. Yoon, S. Byun, S. Dey, and K. Jung, "Speech emotion recognition using multi-hop attention mechanism," in *2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP 2019)*. IEEE, 2019, pp. 2822–2826.

[88] Z. Yang, X. He, J. Gao, L. Deng, and A. Smola, "Stacked attention networks for image question answering," in *2016 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2016)*, 2016, pp. 21–29.

[89] Y. Wang, A. Sun, J. Han, Y. Liu, and X. Zhu, "Sentiment analysis by capsules," in *2018 World Wide Web Conference (WWW 2018)*. ACM, 2018, p. 1165–1174.

[90] S. Sabour, N. Frosst, and G. E. Hinton, "Dynamic routing between capsules," in *31st Annual Conference on Neural Information Processing Systems (NIPS 2017)*. Curran Associates, Inc., 2017, p. 3859–3869.

[91] M. India, P. Safari, and J. Hernando, "Self multi-head attention for speaker recognition," in *Proceedings of the 20th Annual Conference of the International Speech Communication Association (Interspeech 2019)*. ISCA, 2019, pp. 2822–2826.

[92] C. Wu, F. Wu, S. Ge, T. Qi, Y. Huang, and X. Xie, "Neural news recommendation with multi-head self-attention," in *2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP 2019)*. ACL, 2019, pp. 6389–6394.

[93] Y. Wang, W. Chen, D. Pi, and L. Yue, "Adversarially regularized medication recommendation model with multi-hop memory network," *Knowledge and Information Systems*, vol. 63, no. 1, pp. 125–142, 2021.

[94] M. Cornia, M. Stefanini, L. Baraldi, and R. Cucchiara, "Meshed-memory transformer for image captioning," in *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2020)*, 2020, pp. 10 578–10 587.

[95] J. Chen, Y. Lu, Q. Yu, X. Luo, E. Adeli, Y. Wang, L. Lu, A. L. Yuille, and Y. Zhou, "TransUnet: Transformers make strong encoders for medical image segmentation," *arXiv preprint arXiv:2102.04306*, 2021.

[96] P. Zhong, D. Wang, and C. Miao, "Knowledge-enriched transformer for emotion detection in textual conversations," in *2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP 2019)*. ACL, 2019, pp. 165–176.

[97] Y. Zhou, R. Ji, J. Su, X. Sun, and W. Chen, "Dynamic capsule attention for visual question answering," in *33rd AAAI Conference on Artificial Intelligence (AAAI 2019)*, vol. 33, no. 01. AAAI Press, 2019, pp. 9324–9331.

[98] Y. Wang, A. Sun, M. Huang, and X. Zhu, "Aspect-level sentiment analysis using AS-capsules," in *The World Wide Web Conference*, 2019, pp. 2033–2044.

[99] Z. Dai, Z. Yang, Y. Yang, J. Carbonell, Q. Le, and R. Salakhutdinov, "Transformer-XL: Attentive language models beyond a fixed-length context," in *57th Annual Meeting of the Association for Computational Linguistics (ACL 2019)*. ACL, 2019, pp. 2978–2988.

[100] N. Kitaev, Ł. Kaiser, and A. Levskaya, "Reformer: The efficient Transformer," in *8th International Conference on Learning Representations (ICLR 2020)*, 2020.

[101] S. Wang, B. Li, M. Khabsa, H. Fang, and H. Ma, "Linformer: Self-attention with linear complexity," *arXiv:2006.04768*, 2020.

[102] Z. Wu, Z. Liu, J. Lin, Y. Lin, and S. Han, "Lite transformer with long-short range attention," in *8th International Conference on Learning Representations (ICLR 2020)*, 2020.

[103] Y. Tay, D. Bahri, D. Metzler, D.-C. Juan, Z. Zhao, and C. Zheng, "Synthesizer: Rethinking self-attention for transformer models," in *Proceedings of the 38th International Conference on Machine Learning (ICML 2021)*, vol. 139. PMLR, 2021, pp. 10 183–10 192.

[104] Y. Tay, M. Dehghani, D. Bahri, and D. Metzler, "Efficient transformers: A survey," *arXiv:2009.06732*, 2020.

[105] X. Li, J. Song, L. Gao, X. Liu, W. Huang, X. He, and C. Gan, "Beyond RNNs: Positional self-attention with co-attention for video question answering," in *33rd AAAI Conference on Artificial Intelligence (AAAI 2019)*, vol. 33. AAAI Press, 2019, pp. 8658–8665.

[106] A. W. Yu, D. Dohan, M.-T. Luong, R. Zhao, K. Chen, M. Norouzi, and Q. V. Le, "QANet: Combining local convolution with global self-attention for reading comprehension," in *6th International Conference on Learning Representations (ICLR 2018)*, 2018.

[107] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, "BLEU: a method for automatic evaluation of machine translation," in *40th Annual Meeting of the Association for Computational Linguistics (ACL 2002)*. ACL, 2002, pp. 311–318.

[108] S. Banerjee and A. Lavie, "METEOR: An automatic metric for MT evaluation with improved correlation with human judgments," in *2005 Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*. ACL, 2005, pp. 65–72.

[109] R. Sennrich, "Perplexity minimization for translation model domain adaptation in statistical machine translation," in *13th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2012)*. ACL, 2012, pp. 539–549.

[110] M. Popović and H. Ney, "Word error rates: Decomposition over POS classes and applications for error analysis," in *2nd Workshop on Statistical Machine Translation (WMT 2007)*. ACL, 2007, pp. 48–55.

[111] P. Schwarz, P. Matějka, and J. Černockỳ, "Towards lower error rates in phoneme recognition," in *7th International Conference on Text, Speech and Dialogue (TSD 2004)*, ser. LNCS, vol. 3206. Springer, 2004, pp. 465–472.

[112] D. S. Turaga, Y. Chen, and J. Caviedes, "No reference PSNR estimation for compressed pictures," *Signal Processing: Image Communication*, vol. 19, no. 2, pp. 173–184, 2004.

[113] P. Ndajah, H. Kikuchi, M. Yukawa, H. Watanabe, and S. Muramatsu, "SSIM image quality metric for denoised images," in *3rd WSEAS International Conference on Visualization, Imaging and Simulation (VIS 2010)*. WSEAS, 2010, pp. 53–58.

[114] M. A. Rahman and Y. Wang, "Optimizing intersection-over-union in deep neural networks for image segmentation," in *12th International Symposium on Visual Computing (ISVC 2016)*, ser. LNCS, vol. 10072. Springer, 2016, pp. 234–244.

[115] X. Chen, L. Yao, and Y. Zhang, "Residual attention U-net for automated multi-class segmentation of COVID-19 chest CT images," *arXiv:2004.05645*, 2020.

[116] S. Liu, Y. Chen, K. Liu, and J. Zhao, "Exploiting argument information to improve event detection via supervised attention mechanisms," in *55th Annual Meeting of the Association for Computational Linguistics (ACL 2017)*. ACL, 2017, pp. 1789–1798.

[117] C. Liu, J. Mao, F. Sha, and A. Yuille, "Attention correctness in neural image captioning," in *31st AAAI Conference on Artificial Intelligence (AAAI 2017)*. AAAI Press, 2017, pp. 4176–4182.

[118] A. Das, H. Agrawal, L. Zitnick, D. Parikh, and D. Batra, "Human attention in visual question answering: Do humans and deep networks look at the same regions?" *Computer Vision and Image Understanding*, vol. 163, pp. 90 – 100, 2017.

[119] Y. Yu, J. Choi, Y. Kim, K. Yoo, S.-H. Lee, and G. Kim, "Supervising neural attention models for video captioning by human gaze data," in *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2017)*. IEEE Computer Society, 2017.

[120] S. Jain and B. C. Wallace, "Attention is not explanation," in *2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT 2019)*. ACL, 2019, pp. 3543–3556.

[121] S. Wiegreffe and Y. Pinter, "Attention is not not explanation," in *2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP 2019)*. ACL, 2019, pp. 11–20.

[122] A. K. Mohankumar, P. Nema, S. Narasimhan, M. M. Khapra, B. V. Srinivasan, and B. Ravindran, "Towards transparent and explainable attention models," in *58th Annual Meeting of the Association for Computational Linguistics (ACL 2020)*. ACL, 2020, pp. 4206–4216.

[123] K. K. Thekumparampil, C. Wang, S. Oh, and L.-J. Li, "Attention-based graph neural network for semi-supervised learning," *arXiv:1803.03735*, 2018.

[124] D. Nie, Y. Gao, L. Wang, and D. Shen, "ASDNet: Attention based semi-supervised deep networks for medical image segmentation," in *21st International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI 2018)*, ser. LNCS, vol. 11073. Springer, 2018, pp. 370–378.

[125] Y. Alami Mejjati, C. Richardt, J. Tompkin, D. Cosker, and K. I. Kim, "Unsupervised attention-guided image-to-image translation," in *32nd Annual Conference on Neural Information Processing Systems (NIPS 2018)*. Curran Associates, Inc., 2018, pp. 3693–3703.

[126] R. He, W. S. Lee, H. T. Ng, and D. Dahlmeier, "An unsupervised neural attention model for aspect extraction," in *55th Annual Meeting of the Association for Computational Linguistics (ACL 2017)*. ACL, 2017, pp. 388–397.

Gianni Brauwers was born in Spijkenisse, the Netherlands, in 1998. He received the B.S. degree in econometrics and operations research from Erasmus University Rotterdam, Rotterdam, the Netherlands, in 2019, and is currently pursuing the M.S. degree in econometrics and management science at Erasmus University Rotterdam.

He is a Research Assistant at Erasmus University Rotterdam, focusing his research on neural attention models and sentiment analysis.

Flavius Frasincar was born in Bucharest, Romania, in 1971. He received the M.S. degree in computer science, in 1996, and the M.Phil. degree in computer science, in 1997, from Politehnica University of Bucharest, Bucharest, Romania, and the P.D.Eng. degree in computer science, in 2000, and the Ph.D. degree in computer science, in 2005, from Eindhoven University of Technology, Eindhoven, the Netherlands.

Since 2005, he has been an Assistant Professor in computer science at Erasmus University Rotterdam, Rotterdam, the Netherlands. He has published in numerous conferences and journals in the areas of databases, Web information systems, personalization, machine learning, and the Semantic Web. He is a member of the editorial boards of Decision Support Systems, International Journal of Web Engineering and Technology, and Computational Linguistics in the Netherlands Journal, and co-editor-in-chief of the Journal of Web Engineering. Dr. Frasincar is a member of the Association for Computing Machinery.